

Machine Learning

CS 165B

Prof. Matthew Turk

Monday, April 11, 2016

**T
o
d
a
y**

- Classification (cont.)
Chapters 2-3

Notes

- HW#1 answers posted on Gauchospace today
 - Can go over any questions at tomorrow's discussion sessions
 - Main point was to get you to think through several of the concepts important to machine learning
 - Also probability background
 - $P(X, Y)$ vs. $P(X | Y)$
 - Marginalize $P(X, Y)$ to get $P(X)$
- HW#2
 - Posted by this Friday, due following Friday (April 22)
- CS seminar
 - Today at 3:30pm, CS Conference Room (1132 HFH)
 - Sameer Singh, University of Washington
 - “Interactive Machine Learning for Information Extraction”

Classifier margin and loss function

- True class function $c(x) = \begin{cases} +1 & \text{for positive training examples} \\ -1 & \text{for negative training examples} \end{cases}$
- The **scoring classifier** assigns a **margin** $z(x)$ to each instance x :

$$z(x) = c(x)\hat{s}(x)$$


- Positive if the estimate $\hat{s}(x)$ is correct
- Negative if $\hat{s}(x)$ is incorrect
 - Since $\hat{s} > 0$ indicates **positive** estimate and $\hat{s} < 0$ **negative**
- Large **positive margins** mean the classifier is “strongly correct”
- Large **negative margins** are bad – they mean the classifier screwed up!

Classifier margin and loss function

Training data:

Positive class: $c = +1$

Negative class: $c = -1$



Score $\hat{s}(x)$

True class function $c(x)$

Margin $z(x) = c(x)\hat{s}(x)$

How should each training data point impact the classifier learned from this data?

The loss function $L(z)$ will determine this

At this point in the **iterative classifier training algorithm**, which training data points are the most important?

Classifier margin and loss function

- True class function $c(x) = \begin{cases} +1 & \text{for positive training examples} \\ -1 & \text{for negative training examples} \end{cases}$
- The **scoring classifier** assigns a **margin** $z(x)$ to each instance x :

$$z(x) = c(x)\hat{s}(x)$$


- Positive if the estimate $\hat{s}(x)$ is correct
- Negative if $\hat{s}(x)$ is incorrect
 - Since $\hat{s} > 0$ indicates **positive** estimate and $\hat{s} < 0$ **negative**
- Large positive margins mean the classifier is “strongly correct”
- Large negative margins are bad – they mean the classifier screwed up!
- In learning a classifier, we’d like to **penalize** *negative* margins by the use of a **loss function** $L(z)$ that **maps the margin to an associated loss**

$$L : \mathbb{R} \rightarrow [0, \infty)$$

The loss function, $L(z)$

What should the loss function look like?


$$L : \mathbb{R} \rightarrow [0, \infty)$$



*Penalize wrong
classifications more*


The loss function, $L(z)$

- Characteristics of the loss function:
 - For an example on the decision boundary, $L(0) = 1$
 - $L(z) \geq 1$ for $z < 0$
 - $0 \leq L(z) < 1$ for $z > 0$



The loss function, $L(z)$

- Loss functions are often used in **optimization** problems (to minimize a function) that lead to modifying **weights** in training
 - Typically it is squared – thus the mapping to $[0, \infty)$
- To help make this solvable, the loss function is often chosen to be **convex**, since optimizing a convex function is computationally more tractable



A **convex function** lies below the line connecting any two points on the function

Ranking classifier

- The scores from a **scoring classifier** may not be particularly meaningful – they are not derived from any “true” scores – so it may be preferable to ignore the **magnitude** and just keep the **order** of the scores on a set of instances
 - This is less sensitive to **outliers** – i.e., more robust to noise/errors
- All positive examples should (ideally) be ranked higher than all negative examples
 - Exceptions to this are **ranking errors**
 - Count the ranking errors (*err*): For all (**pos**, **neg**) example pairs, how many rank **neg** higher than **pos**?
 - Ties count 1/2

Ranking error rate: $rank\text{-}err = err / pN$

Ranking accuracy: $rank\text{-}acc = 1 - rank\text{-}err$

Ranking classifier performance

Score: Low  High

of ranking errors? Ranking error rate? Ranking accuracy?

2 $2/(5)(5) = 0.08$ 0.92

Low  High

of ranking errors? Ranking error rate? Ranking accuracy?


9 $9/(8)(12) = 0.09375$ 0.90625

Low  High

of ranking errors? Ranking error rate? Ranking accuracy?


9 $9/(8)(12) = 0.09375$ 0.90625

Ranking classifier and the coverage curve



Move the **decision line** and count:

FP = ? TP = ?



This is the coverage curve


If we normalize to a square graph, we get the ROC curve

The area under the ROC curve is the **ranking accuracy**

Ranking classifier and the coverage curve

Low ● ● ● ● ● ● ● ● ● ● High

What about this case? Looks like the ranking is terrible!



What is the **ranking accuracy**? Zero (0%)

Q: What would the ranking accuracy be of this ranking?

Low ● ● ● ● ● ● ● ● ● ● High




One (100%)

Classifier design – operating point



- You, as a classifier designer, can often move decision boundaries (modify thresholds) to make the **false positive rate** as high or as low as you wish
 - A very high threshold (don't let anything through!) results in no **false positives** – but lots of **false negatives**
 - A very low threshold (let everything through!) results in no **false negatives** – but lots of **false positives**
- This doesn't necessarily make the classifier better or worse – it just changes the **operating point** of the classifier
- This is often application-specific:
 - When might false positives be especially undesirable?
 - When might false negatives be especially undesirable?
 - We can encode these preferences in a **cost function** to compute an optimal threshold, given this information

Classifier design

Which classifier is best: A, B, or C?



A or B?





It depends on what you want to optimize:



- TPR, FPR, error rate, accuracy, precision, ...

Classifier design

A or B?



A or B?




It also depends on how good your **features** are!
– Either *raw* features or *constructed* features

Feature separation vs. classifier design

Placing the separating boundary = **classifier design**

Increasing the feature separation = **feature construction**



We can design a better classifier starting with the **features** in **C**!

Class probability estimation

A **class probability estimator** is a scoring classifier that outputs probabilities over the k classes – i.e., a mapping:

$$\hat{p} : \mathcal{X} \rightarrow [0, 1]^k$$

where

$$\sum_{i=1}^k \hat{p}_i(x) = 1$$

A key issue here is that we generally do not have access to the **true probabilities** for training data.

- E.g., an email is either spam or ham – it doesn't have a probability of being spam!
- So how can we train to learn such probabilities?

Empirical probabilities

- In machine learning, we often calculate *empirical probabilities*
 - i.e., calculate relative frequencies from the available data

N_i instances of k classes C_i in the training data S :

$$\text{Relative frequency} = \frac{N_i}{|S|} = \hat{p}_i$$

- But this can be problematic, especially with small amounts of training data
 - Probabilities of 0 and 1 generally should be avoided
- There are various common ways to smooth or correct the relative frequencies to avoid 0 and 1
 - E.g., Laplace correction and m-estimate:

Add a pseudo-count to each class

$$\text{Laplace correction} = \frac{N_i + 1}{|S| + k}$$

Choose number of pseudo-counts m and their class distribution π_i

$$\text{m-estimate} = \frac{N_i + m\pi_i}{|S| + m} \quad \sum_i \pi_i = 1$$

Empirical probabilities

Training data set S

C_1 : 7 instances

C_2 : 14

C_3 : 0

C_4 : 4

$$\text{Relative frequency} = \frac{N_i}{|S|} = \hat{p}_i$$

$$\begin{aligned} |S| &= 25 & \hat{p}_1 &: 7/25 = 0.28 \\ & & \hat{p}_2 &: 14/25 = 0.56 \\ & & \hat{p}_3 &: 0/25 = 0.0 \\ & & \hat{p}_4 &: 4/25 = 0.16 \end{aligned}$$

$$\text{m-estimate} = \frac{N_i + m\pi_i}{|S| + m} \quad \Sigma_i \pi_i = 1$$

$$m = 40 : \{10, 10, 10, 10\}$$

$$m = 40 : \{10, 0, 18, 12\}$$

$$\hat{p}_1: 17/65 = 0.26$$

$$\hat{p}_2: 24/65 = 0.37$$

$$\hat{p}_3: 10/65 = 0.15$$

$$\hat{p}_4: 14/65 = 0.22$$

$$\hat{p}_1: 17/65 = 0.26$$

$$\hat{p}_2: 14/65 = 0.22$$

$$\hat{p}_3: 18/65 = 0.28$$

$$\hat{p}_4: 16/65 = 0.25$$

$$\text{Laplace correction} = \frac{N_i + 1}{|S| + k}$$

$$\hat{p}_1: 8/29 = 0.28$$

$$\hat{p}_2: 15/29 = 0.52$$

$$\hat{p}_3: 1/29 = 0.03$$

$$\hat{p}_4: 5/29 = 0.17$$

Multi-class classification

- Many classification problems involve **multiple classes**
- Performance can be described with the **multi-class contingency table**
 - Not including the marginals, also known as the **confusion matrix**
 - We can compute **accuracy**, per-class **precision**, per-class **recall**...

	<i>Predicted</i>			
<i>Actual</i>	15	2	3	20
	7	15	8	30
	2	3	45	50
	24	20	56	100

$$\text{Accuracy} = (15+15+45)/100 = 0.75$$

$$\text{Class 1 precision} = 15/24 = 0.63$$

$$\text{Class 1 recall} = 15/20 = 0.75$$

Etc.