

# Machine Learning

CS 165B

Prof. Matthew Turk

Wednesday, April 13, 2016

**T**  
**o**  
**d**  
**a**  
**y**

- Classification (cont.)  
Chapters 2-3
- Concept learning  
Chapter 4

# Notes

---

- HW#2
  - Posted on Friday, due next Friday (April 22)

# Multi-class classification

---

- Many classification problems involve **multiple classes**
- Performance can be described with the **multi-class contingency table**
  - Not including the marginals, also known as the **confusion matrix**
  - We can compute **accuracy**, per-class **precision**, per-class **recall**...

		<i>Predicted</i>			
<i>Actual</i>	15	2	3	20	
	7	15	8	30	
	2	3	45	50	
	24	20	56	100	

$$\text{Accuracy} = (15+15+45)/100 = 0.75$$


$$\text{Class 1 precision} = 15/24 = 0.63$$

$$\text{Class 1 recall} = 15/20 = 0.75$$


Etc.

# K-class classifiers

- How to build a *k*-class classifier?
  - We can combine several binary classifiers, e.g.:
    - **One-versus-rest scheme** – learn  $k-1$  models, apply in sequence
      - C1 vs. { C2, C3, C4 }
      - C2 vs. { C3, C4 }
      - C3 vs. C4




- **One-versus-rest scheme** – learn a *one-class* model for each class
  - C1 vs. { C2, C3, C4 }
  - C2 vs. { C1, C3, C4 }
  - C3 vs. { C1, C2, C4 }
  - C4 vs. { C1, C2, C3 }




# K-class classifiers

- **One-versus-one scheme** – learn a model for each pair of classes
  - Train  $k(k-1)/2$  binary classifiers, apply them all to  $x$  and **vote**




- **One-versus-one scheme** with a decision tree:



# Example: A 3-class linear classifier

---

Classify instances into three classes  $\{A, B, C\}$  using three **linear discriminant functions** classifying (A vs. B), (A vs. C), and (B vs. C)



# Regression – another predictive ML task


---

- In the classification tasks we've been discussing, the **label space** was a discrete set of classes
  - Classification, scoring, ranking, probability estimation
- **Regression** learns a function (the **regressor**) that is a mapping  $\hat{f} : \mathcal{X} \rightarrow \mathbb{R}$  from examples –  $f(x_i)$ 
  - I.e., the **target variable** (output) is real-valued
- Assumption: the examples will be noisy, so watch out for **overfitting** – need to capture the general trend or shape of the function, not exactly match every data point
  - E.g., if fitting an **N-degree polynomial** to the training data (thus  $N+1$  parameters to estimate), choose one as low degree as possible
- The number of data points should be much greater than the number of parameters to be estimated!
  - How much data is needed? An open question in ML....

# Regression example

Training data

$x$	$f(x)$
1.0	1.2
2.5	2.0
4.1	3.7
6.1	4.6
7.9	7.0



— Piecewise linear fit

- - - Globally linear fit


The regression function **may or may not** fit the training data exactly



# Regression example

Training data

$x$	$f(x)$
1.0	1.2
2.5	2.0
4.1	3.7
6.1	4.6
7.9	7.0




{1, 2, 3, 4}-order functions

# Regression

---

- We'll generally estimate a regression function based on some function of the *residual*, the different between the *estimate* and the *label* (the true value):

$$r(x) = f(x) - \hat{f}(x)$$



True function                  Regression function

- That function is (again) the *loss function L*
  - The most common loss function for regression is the *squared residual*:

$$L(x) = r^2(x) = (f(x) - \hat{f}(x))^2$$

- However, this is sensitive to *outliers* (large errors have a disproportionately large effect), so often a function that minimizes large errors is used – the result is called a *robust estimator*

# Concept Learning

Chapter 4 in the textbook

*Logical Models: tree models and rule models*

# Concept learning

---

- Concept learning means learning (typically **binary**) concepts from examples
  - The learned concept is the **positive** class
  - Everything else is the **negative** class
- We'll now use **logical models** – logical expressions describe concepts and divide the **instance space** appropriately
  - See “Background 4.1” on p. 105 in the textbook (or take CS 40 or CS 165A!) for an overview of the logical concepts and notation
  - Propositional logic
    - Logical manipulation of propositions (symbols that have values)
  - (First-order) predicate logic
    - Add variables, predicates (binary functions), functions, and variable quantification (**for all x..., there exists an x such that...**)

# Propositional (Boolean) logic

---

- Symbols represent **propositions** (statements of fact, sentences)
  - $P$  means “San Francisco is the capital of California”
  - $Q$  means “It is raining in Seattle”
  - $Length = 3$  means “The value of the feature  $Length$  is 3”
  - $Teeth = many$  means “The value of the feature  $Teeth$  is  $many$ ”
- Expressions are generated by combining proposition symbols with Boolean (logical) **connectives**
  - $True, false$ , propositional symbols
  - $feature/value$  relations – e.g.,  $feature = value$ ,  $feature < value$ , ...
  - $( )$ ,  $\neg$  (not),  $\wedge$  (and),  $\vee$  (or),  $\Rightarrow$  (implies),  $\Leftrightarrow$  (equivalent)

# Propositional logic

---

- Semantics
  - Defined by clearly interpreted symbols and straightforward application of truth tables
  - Rules for evaluating truth: Boolean algebra
  - Simple method: truth tables

$P$	$Q$	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>
<i>True</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>
<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>

**$2^N$  rows for  $N$  propositions**

# Concept learning

---

- In **concept learning**, we want to learn a **Boolean function** over a set of attributes+values
  - I.e., derive a Boolean function from training examples
    - Positive and negative examples of the concept
      - Positive:  $\text{Temperature} = \text{high} \wedge \text{Coughing} = \text{yes} \wedge \text{Spots} = \text{yes}$
      - Negative:  $\text{Temperature} = \text{medium} \wedge \text{Coughing} = \text{no} \wedge \text{Spots} = \text{yes}$
    - This is our **hypothesis**
- The **target concept**  $c$  is the true concept
  - We want the hypothesis to be a good estimate of the true concept
  - Thus we wish to find  $h$  (or  $\hat{c}$ ) such that  $h = c$  (or  $\hat{c} = c$ )
- The hypothesis is a **Boolean function over the features**
  - E.g., some combinations of  $\{\text{Temperature}, \text{Coughing}, \text{Spots}\}$  are in the concept, and others are not in the concept

# The hypothesis space

---

- Using a set of features, what concepts can possibly be learned?
- The space of all possible concepts is called the **hypothesis space**
  - What is the hypothesis space for a given problem?
- First, how many possible **instances** are there for a given set of **features**?
  - In set theory, the Cartesian product of all the features
  - $F_1 \times F_2 \times \dots \times F_N$
  - All combinations of feature values
  - UCSB courses: Quarter (4), Dept (40), courselevel (2), topic (500)
  - $4 \times 40 \times 2 \times 500 = 160,000$  possible instances
    - E.g., (spring, CS, ugrad, ML), (fall, Music, grad, StringTheory), ...



# The hypothesis space

---

- The hypothesis space is the number of binary functions on these instances, which is...  $2^{160,000}!!$ 
  - I.e., the number of different sets you can make from 160,000 elements
  - Or if you laid out all possible instances, the number of different contours you could draw separating some instances from the rest
  - Each of these hypotheses... sets... contours... defines a concept
- The challenge in concept learning is deciding which hypothesis is best, given the training data
  - As with all problems in machine learning, generalization is of key importance – we don't only want to memorize the training data (the overfitting problem)
  - We want to learn a concept that will generalize well to new, unseen instances

# The conjunctive hypothesis space

---

- Let's limit our hypothesis space to **conjunctive** concepts – i.e., hypotheses that can be expressed as **a conjunction of literals** (features)

Quarter=?  $\wedge$  Dept=?  $\wedge$  courselevel=?  $\wedge$  topic=?

- We add “**absence**” or “**don't care**” to each feature, so now the total number of combinations is  $5 \times 41 \times 3 \times 501 = 308,115$ 
  - That's a lot, but much better than  $2^{160,000}$ ! (between  $2^{18}$  and  $2^{19}$ )
- The most general hypothesis is (X, X, X, X), which includes all possible instances
  - (fall, X, X, X) is the **concept** of all fall quarter courses
  - (fall, CS, grad, X) is the **concept** of all CS graduate courses in the fall
- In this conjunctive hypothesis space, we can't represent concepts like “**all courses in AI or Graphics**”

# An example hypothesis space

**Target concept:**  
*c = Pays income taxes*

Two binary features:

		UCSB student	
		Yes	No
Age < 21	Yes	A	B
	No	C	D

Instance space:

$\{\text{Age} \times \text{Student}\}$

$2 \times 2 = 4$  instances:

(Yes, Yes) – A      (Yes, No) – B  
 (No, Yes) – C      (No, No) – D

How many possible hypotheses are there?

$2^4 = 16$  possible hypotheses (concepts)

The training example (Yes, No) provides evidence for which hypotheses?

– All the ones that contain B

Now what if we observe a second training example (No, No)?

{A}	{A, D}
★ {B}	★ {B, C}
{C}	★ {B, C, D}
{D}	{A, C, D}
★ {A, B}	★ {A, B, D}
{C, D}	★ {A, B, C}
{A, C}	★ {A, B, C, D}
★ {B, D}	{ } or $\phi$

# Our example using conjunctive hypothesis space

		UCSB student	
		Yes	No
Age < 21	Yes	<b>A</b>	<b>B</b>
	No	<b>C</b>	<b>D</b>

Instance space:  
 $\{\text{Age} \times \text{Student}\}$

CHS: Hypotheses that can be represented as  
 $\text{Age} = \{\text{Yes}, \text{No}, \text{X}\} \wedge \text{Student} = \{\text{Yes}, \text{No}, \text{X}\}$

That's 9 hypotheses:

(Yes, Yes)	(Yes, No)	(Yes, X)
(No, Yes)	(No, No)	(No, X)
(X, Yes)	(X, No)	(X, X)

$\{A\}$	$\{A, D\}$
$\{B\}$	$\{B, C\}$
$\{C\}$	$\{B, C, D\}$
$\{D\}$	$\{A, C, D\}$
$\{A, B\}$	$\{A, B, D\}$
$\{C, D\}$	$\{A, B, C\}$
$\{A, C\}$	$\{A, B, C, D\}$
$\{B, D\}$	$\{\}$ or $\phi$



Conjunctive = combining rows and columns via AND (not by OR)

# An example of CHS learning

---

Suppose you come across a number of sea animals that you suspect belong to the same species. You observe their **length** in meters, whether they have **gills**, whether they have a prominent **beak**, and whether they have few or many **teeth**. The first animal can be described by the following **conjunction** of features:

$\text{Length} = 3 \wedge \text{Gills} = \text{no} \wedge \text{Beak} = \text{yes} \wedge \text{Teeth} = \text{many}$

The next one has the same characteristics but is a meter longer, so you drop the length condition and generalize the conjunction to

$\text{Gills} = \text{no} \wedge \text{Beak} = \text{yes} \wedge \text{Teeth} = \text{many}$

The third animal is again 3 meters long, has a beak, no gills and few teeth, so your description becomes

$\text{Gills} = \text{no} \wedge \text{Beak} = \text{yes}$

All remaining animals satisfy this conjunction, and so your hypothesis is formed.

Someone tells you what these animals are called: **Dolphins**

# An example of CHS learning

---

We took a **specific-to-general** approach in coming up with a hypothesis here.

Instances:

Hypotheses:

(1)  $\text{Length} = 3 \wedge \text{Gills} = \text{no} \wedge \text{Beak} = \text{yes} \wedge \text{Teeth} = \text{many}$

$\text{Length} = 3 \wedge \text{Gills} = \text{no} \wedge \text{Beak} = \text{yes} \wedge \text{Teeth} = \text{many}$

(2)  $\text{Length} = 4 \wedge \text{Gills} = \text{no} \wedge \text{Beak} = \text{yes} \wedge \text{Teeth} = \text{many}$

$\text{Length} = \text{X} \wedge \text{Gills} = \text{no} \wedge \text{Beak} = \text{yes} \wedge \text{Teeth} = \text{many}$

(3)  $\text{Length} = 3 \wedge \text{Gills} = \text{no} \wedge \text{Beak} = \text{yes} \wedge \text{Teeth} = \text{few}$

$\text{Length} = \text{X} \wedge \text{Gills} = \text{no} \wedge \text{Beak} = \text{yes} \wedge \text{Teeth} = \text{X}$

# An example of CHS learning

---

Features and possible values:

Length = { 3, 4, 5 }

Gills = { yes, no }

Beak = { yes, no }

Teeth = { few, many }

In this problem, there are  $3 \times 2 \times 2 \times 2 = 24$  possible **instances** and  $2^{24}$  possible **hypotheses** over the instances (about 16.8 million)


But with the **conjunctive hypothesis space**, we have only  $4 \times 3 \times 3 \times 3 = 108$  possible conjunctive hypotheses

- In our earlier example, we went from 16 hypotheses to 9 using CHS
- Here we go from 16.8 million to 108

Let's visualize the **conjunctive hypothesis space** for this problem:

# The Conjunctive Hypothesis Space

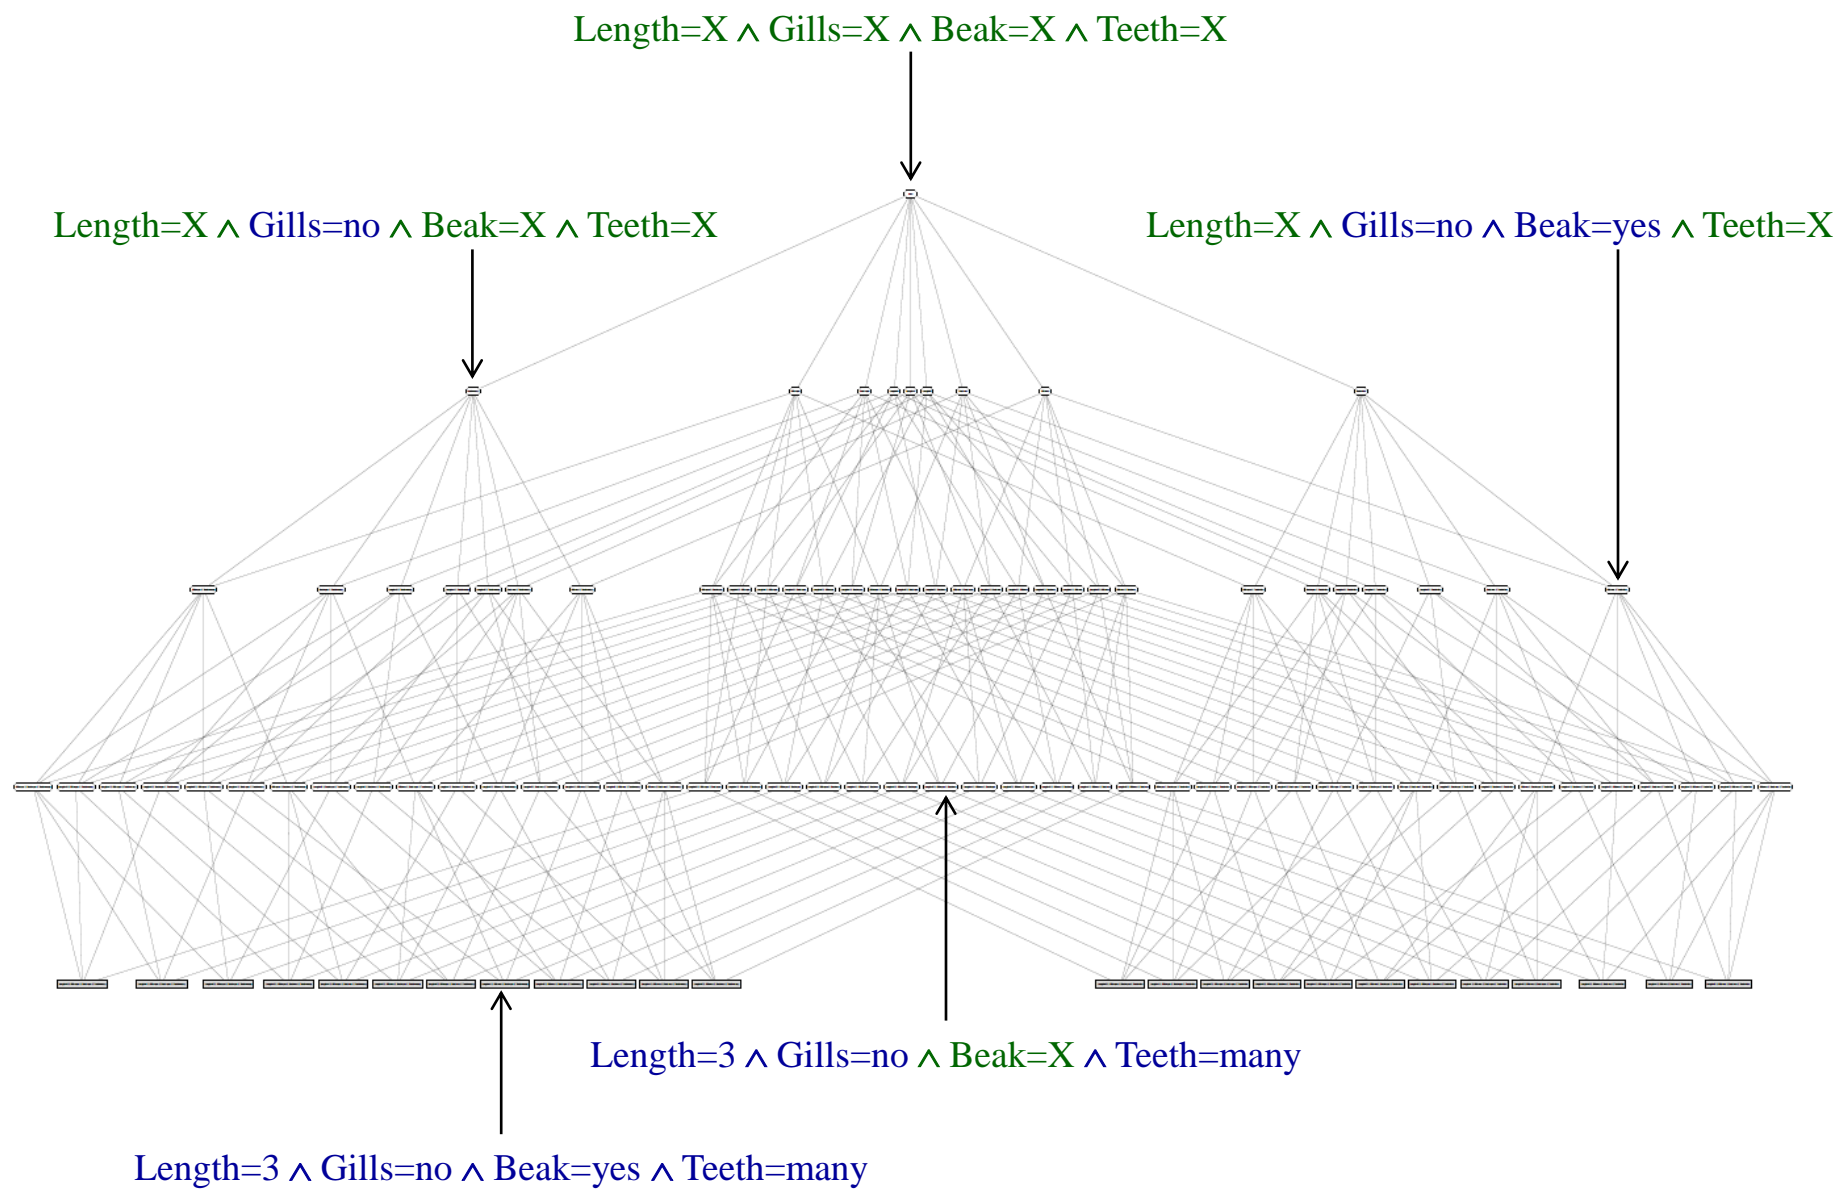
108 nodes (conjunctive concepts)



This connects upward to every **more general** hypothesis that includes it



# The Conjunctive Hypothesis Space




Number of “Don’t know” (X) increases by 1 each level

# The Conjunctive Hypothesis Space

Without writing the “Don’t know”s

Note that  $\wedge \text{True}$  can be appended to any proposition, so this is what’s left when all the literals are removed



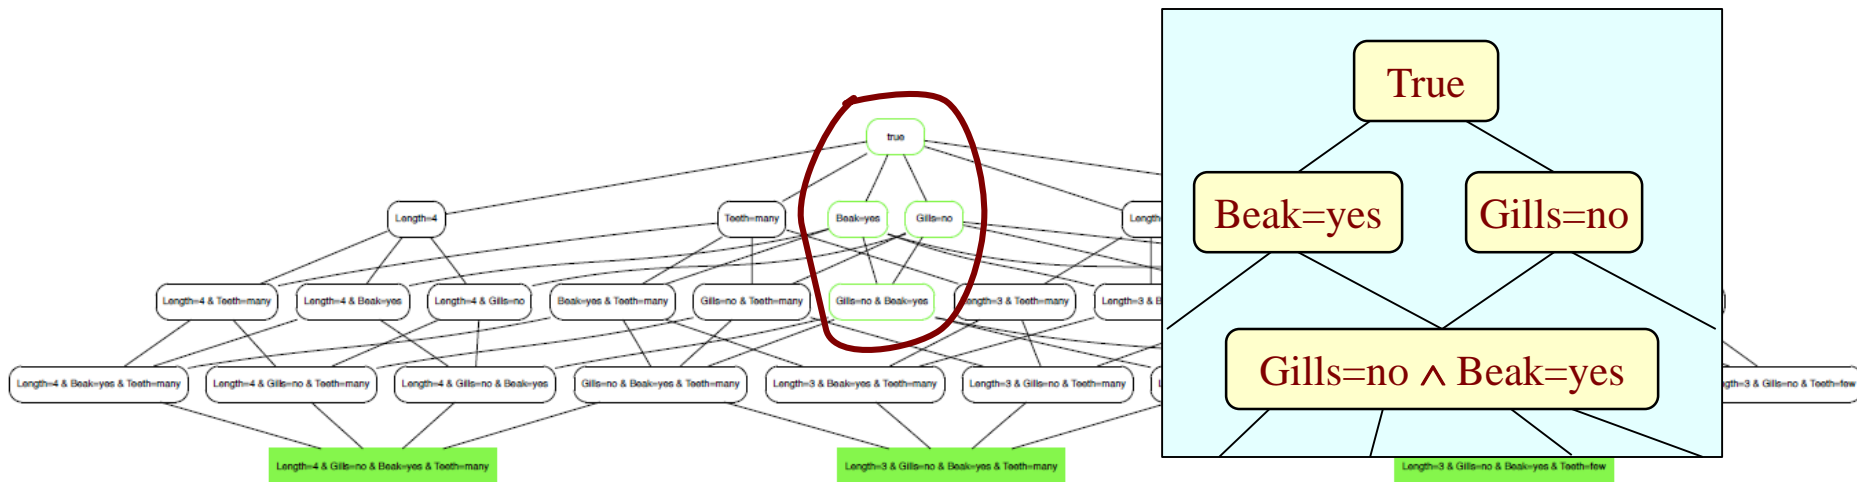
# Reducing the hypothesis space

---

- Given that hypothesis space and our training examples:
  - Length = 3  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = many
  - Length = 4  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = many
  - Length = 3  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = few
- Let's rule out all the hypotheses (concepts) that don't include **at least one** of the instances in our example
  - I.e., **delete nodes** that don't fit with at least one training example
  - This leaves us with just **32** conjunctive concepts (out of the original **108**)

# Reducing the hypothesis space

- But if we require hypotheses to cover all three examples, we're left with only **four** concepts



- Let's choose the **least general** of these as our result – the concept defined by our training data

**Gills = no ∧ Beak = yes**

# Least general generalization (LGG)

---

- We want to **generalize** beyond our specific training data, but not too much – the most general hypothesis is to accept everything
- Thus we'd like the *least general generalization*
  - General enough to include all of our training data, but no more general than that
- Referring to our classification terminology, the more general our hypothesis, the lower our...
  - False negative rate     $H(x) = \text{True}$
- And the less general (more specific) our hypothesis, the lower our...
  - False positive rate     $H(x) = \text{False}$
- So, our approach to generating a hypothesis/concept should depend on the needs of our application