

Machine Learning

CS 165B

Prof. Matthew Turk

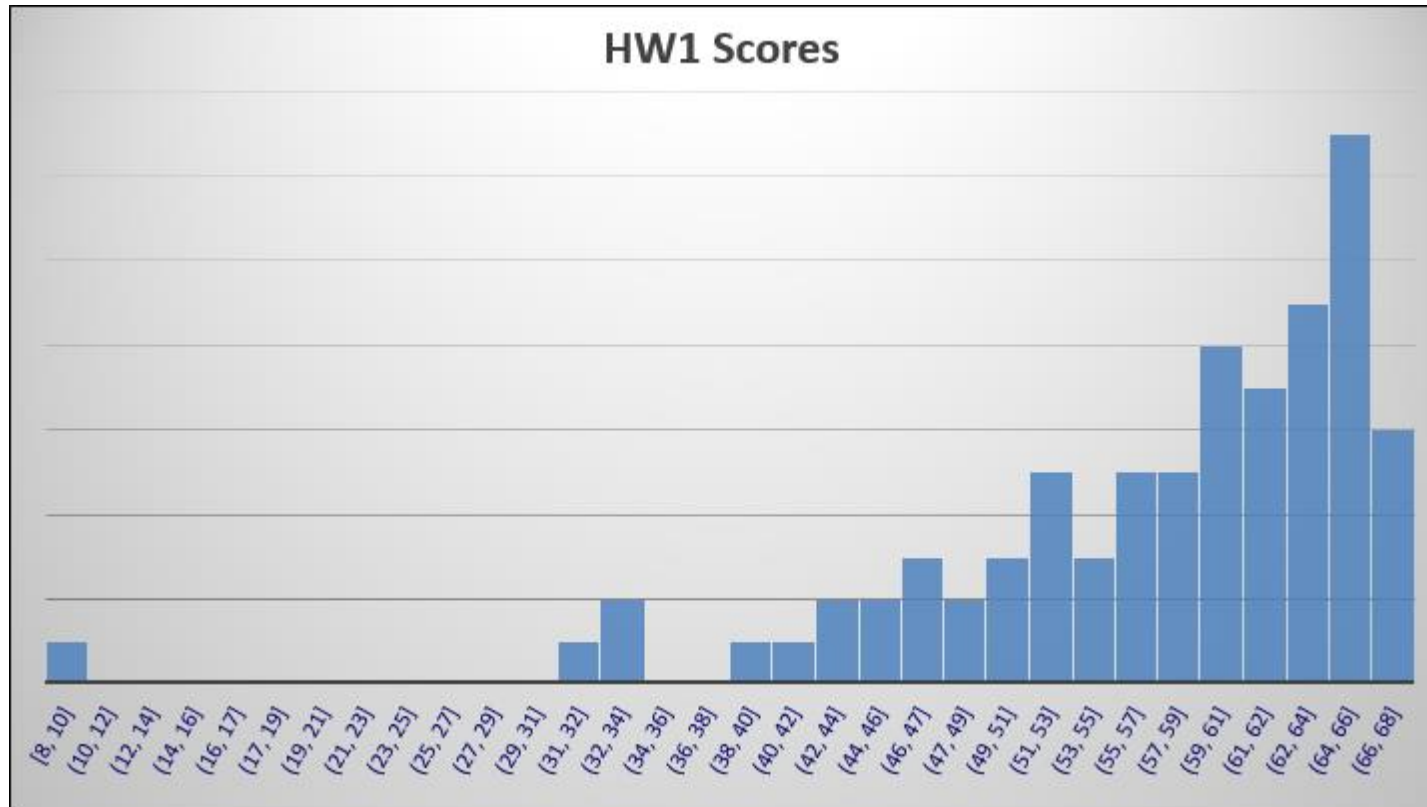
Wednesday, April 20, 2016

T
o
d
a
y

- Decision trees
- Linear learning models (Ch 7)

Notes

- HW#1 scores
 - Ave = $56.5/68 = 83\%$
 - Median = $59/68 = 87\%$

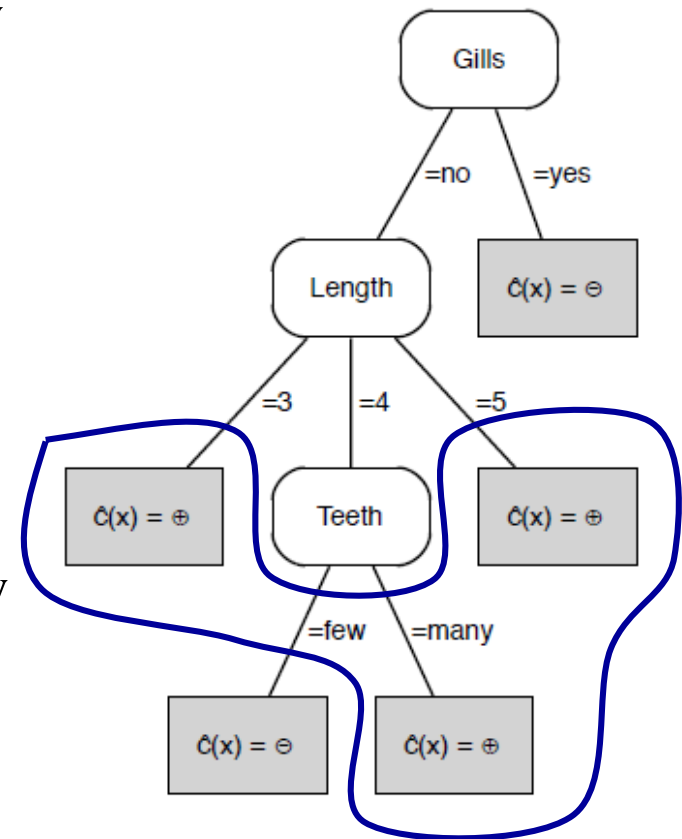


Notes

- HW#2 due 4:30pm Friday
- Midterm – Monday, May 2, in class
 - Covers material through next Wednesday
 - Brief review in class next Wednesday
 - Practice midterm will be supplied
 - Closed book/notes
 - Exception: You may bring one 8.5”x11” sheet of paper with your notes (both sides)
 - I’ll also provide some information, formulas, etc. (will be included with the practice midterm)
- Gauchospace expiration notifications...

Decision trees

- A **decision tree** partitions the **instance space** by branching on feature values (**literals**), with **leaves** representing the learned concept
- Each leaf represents a **conjunction** of literals on its path
- The learned concept is the **disjunction** of the positive leaves
 - $L_1 \vee L_2 \vee L_3 \vee \dots$
- Decision trees are maximally expressive – they can separate any consistently labeled data
 - Thus more powerful than the **conjunctive hypothesis space** we just discussed



Ideally, each leaf contains only positives or only negatives from the training data

Key question: Which features (and in what order) will accomplish this best?

Decision trees

Note: This is very similar to ID3, a well-known DT algorithm – in yesterday's discussion session

- Tree models can be used for **classification**, **ranking**, **probability estimation**, **regression**, and **clustering**
- Recursive generic tree learning procedure:

Algorithm $\text{GrowTree}(D, F)$ – grow a feature tree from training data.

Input : data D ; set of features F .

Output : feature tree T with labelled leaves.

if $\text{Homogeneous}(D)$ **then return** $\text{Label}(D)$;

$S \leftarrow \text{BestSplit}(D, F)$; // e.g., BestSplit-Class (Algorithm 5.2)

split D into subsets D_i according to the literals in S ;

for each i **do**

if $D_i \neq \emptyset$ **then** $T_i \leftarrow \text{GrowTree}(D_i, F)$;

else T_i is a leaf labelled with $\text{Label}(D)$;

end

return a tree whose root is labelled with S and whose children are T_i

100%?
99%?
80%?

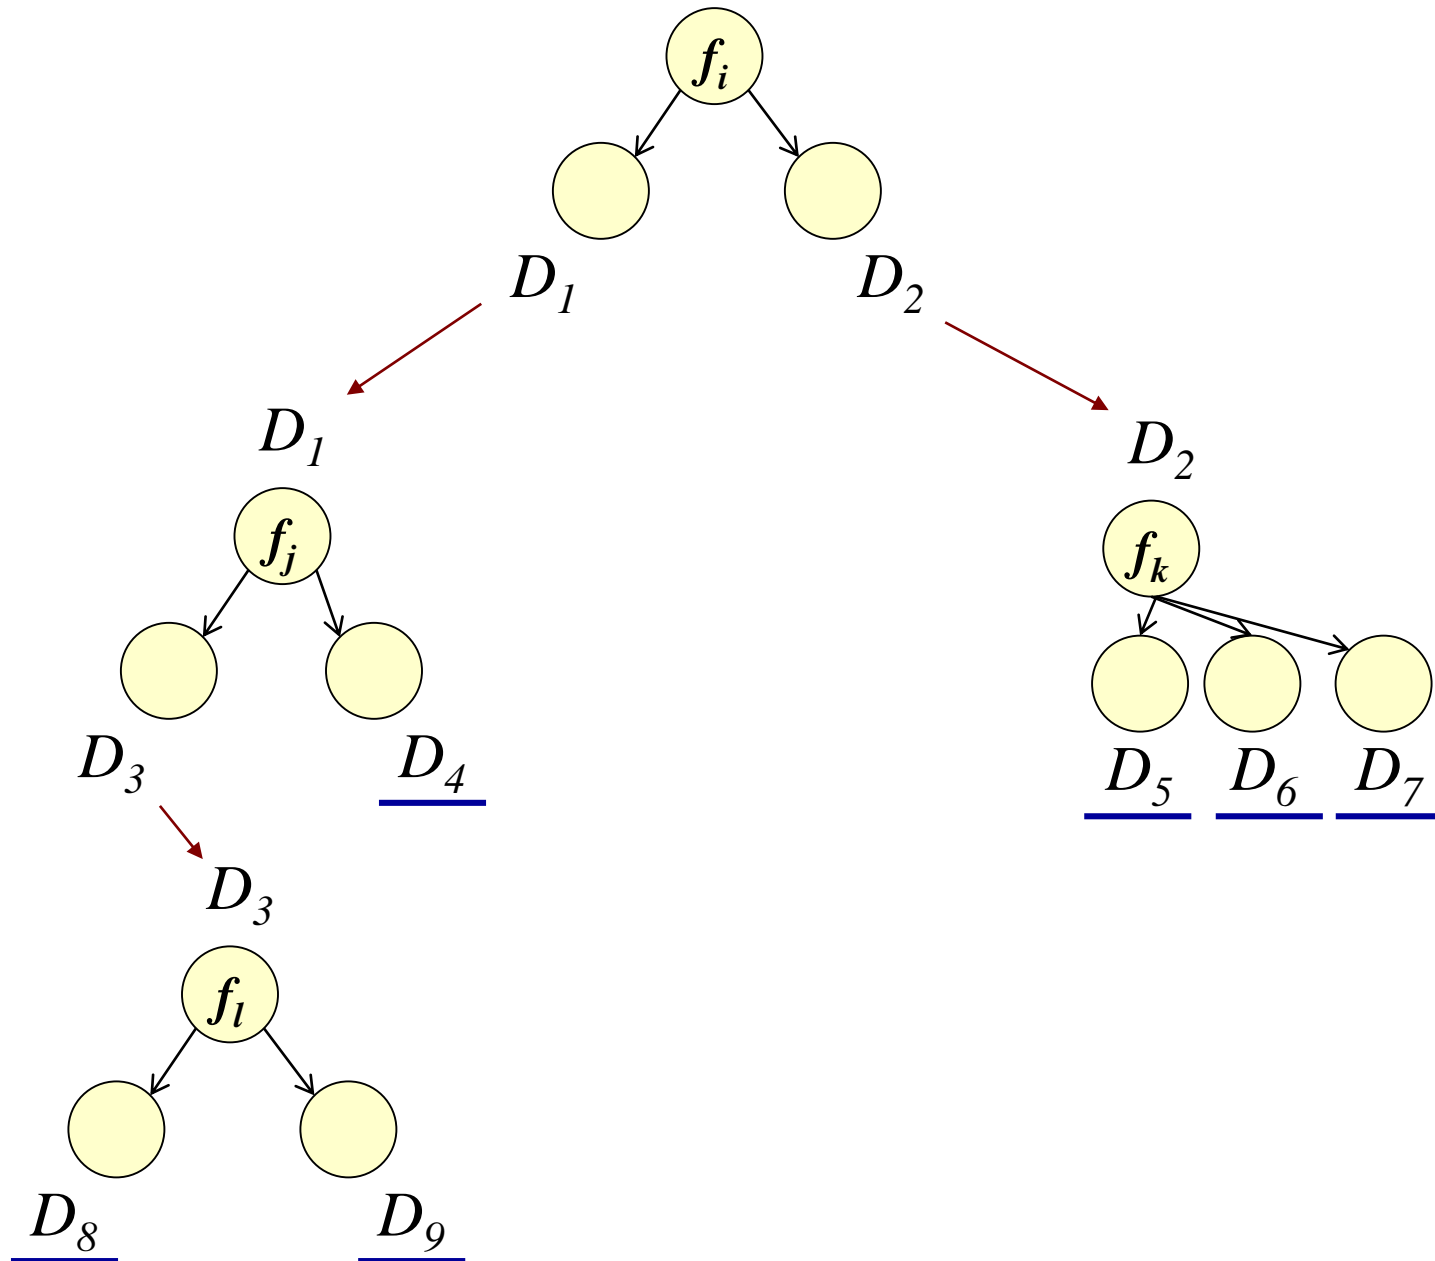
Most useful
feature

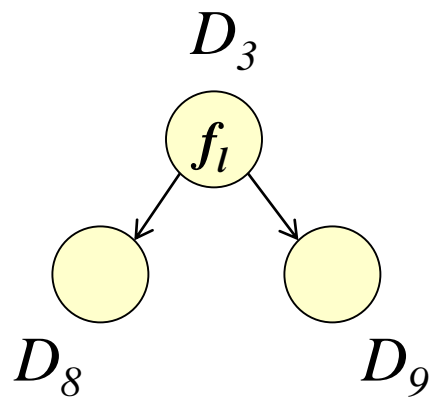
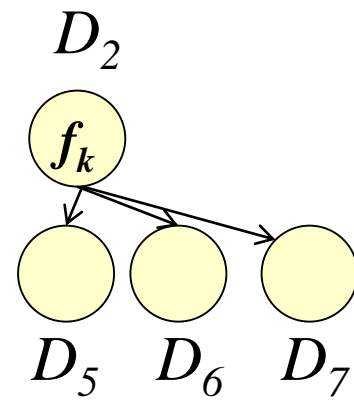
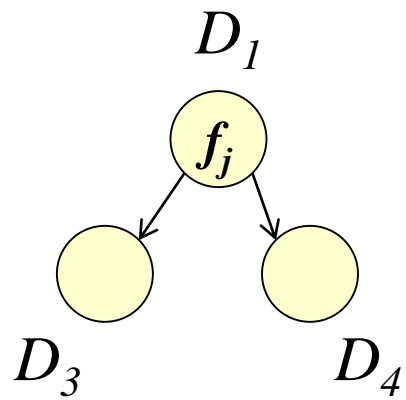
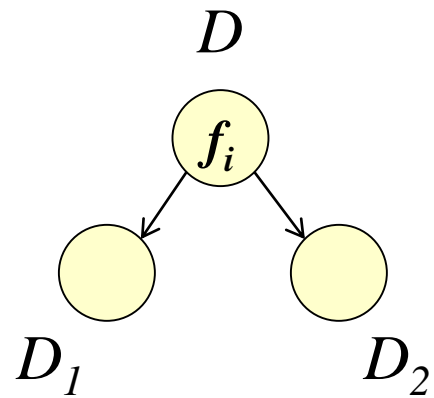


Divide-and-conquer approach: build a tree for each subset of the data, then merge into a single tree

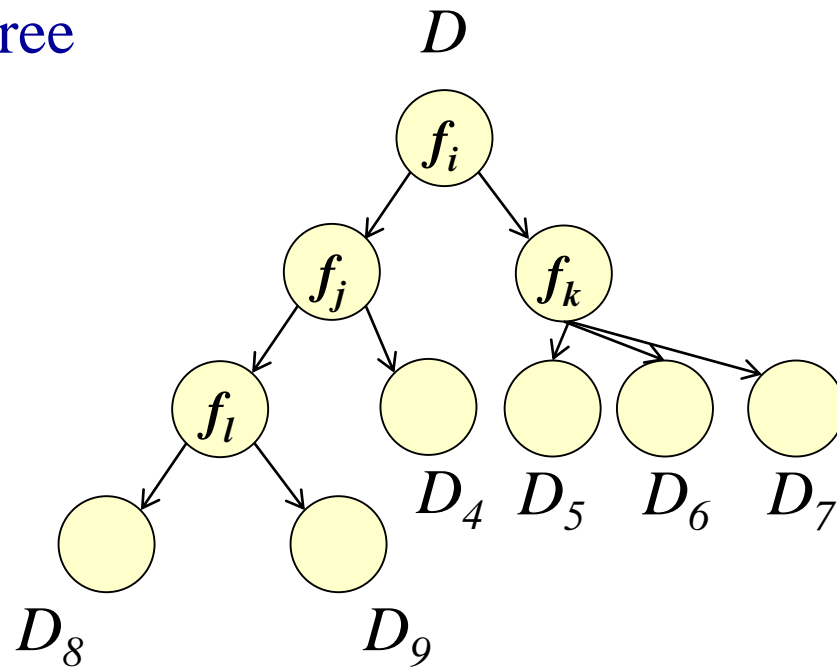
Training Data D

Features $F = \{f_1, f_2, \dots\}$





Feature Tree



Decision trees

Note: This is very similar to ID3, a well-known DT algorithm – in yesterday's discussion session

- Tree models can be used for **classification**, **ranking**, **probability estimation**, **regression**, and **clustering**
- Recursive generic tree learning procedure:

Algorithm $\text{GrowTree}(D, F)$ – grow a feature tree from training data.

Input : data D ; set of features F .

Output : feature tree T with labelled leaves.

if $\text{Homogeneous}(D)$ **then return** $\text{Label}(D)$;

$S \leftarrow \text{BestSplit}(D, F)$; // e.g., BestSplit-Class (Algorithm 5.2)

split D into subsets D_i according to the literals in S ;

for each i **do**

if $D_i \neq \emptyset$ **then** $T_i \leftarrow \text{GrowTree}(D_i, F)$;

else T_i is a leaf labelled with $\text{Label}(D)$;

end

return a tree whose root is labelled with S and whose children are T_i

100%?
99%?
80%?

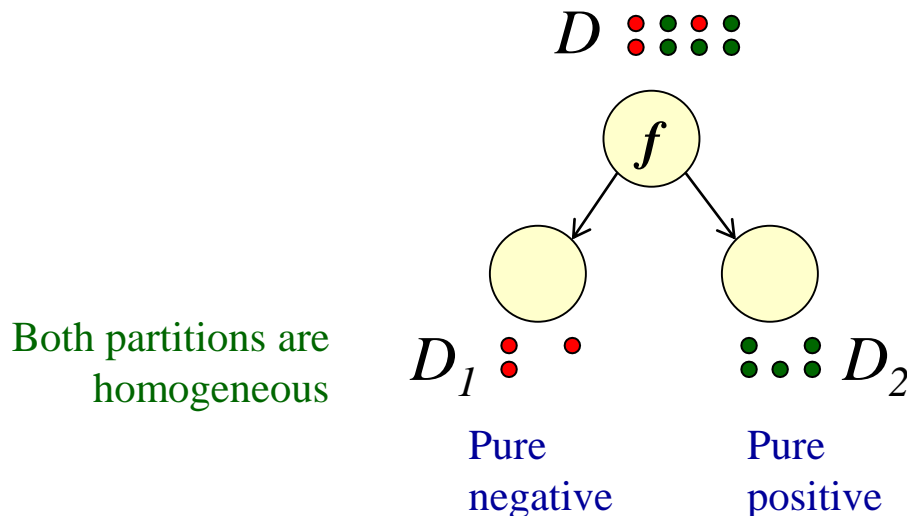
Most useful
feature



Divide-and-conquer approach: build a tree for each subset of the data, then merge into a single tree

BestSplit

- $\text{BestSplit}(D, F)$ – what feature $f \in F$ will produce the best split (partitioning) of the training data $D = \{ D_i \}$?
- What's a good split/partitioning?
 - One that produces **pure** partitions D_i , each of which contains **only instances from a single class**
 - E.g., in a binary classification problem, D_1 has only positive examples and D_2 has only negative examples



Which feature f is best for this?

How to measure partition **purity** if not completely homogeneous?

Impurity

- In the **binary case**, we have P positives and N negatives in the data
 - The best split would be a feature that divides the data D into two **pure** partitions: D_1 with the P positives and D_2 with the N negatives
- So a measure of partition **impurity** should be minimum when the data are 100% positives or negatives, and maximum when 50/50
- Like with empirical probabilities, we can **estimate impurity** by counting. We define **the proportion of positives** in D_i as:

$$\dot{p} = \frac{P}{P + N}$$

- **Impurity** is a function of \dot{p}
 - Should be zero when $\dot{p} = 0$ or 1 , and maximum when $\dot{p} = 0.5$
 - We can write impurity as **Imp(D)** or **Imp(\dot{p})**

Impurity functions

Minority class

$$\text{Imp}(\dot{p}) = \min(\dot{p}, 1-\dot{p})$$

Gini index

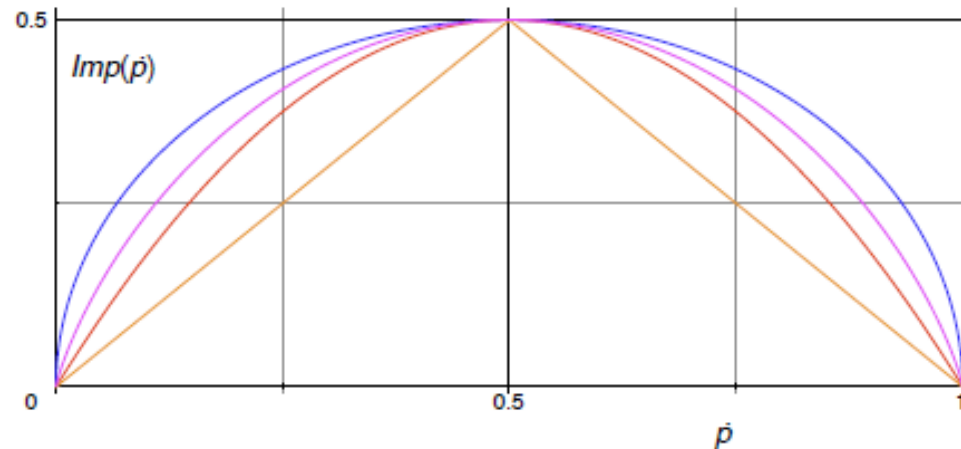
$$\text{Imp}(\dot{p}) = 2\dot{p}(1-\dot{p})$$

Entropy

$$\text{Imp}(\dot{p}) = -\dot{p} \log_2(\dot{p}) - (1-\dot{p}) \log_2(1-\dot{p})$$

$\sqrt{\text{Gini}}$ index

$$\text{Imp}(\dot{p}) = \sqrt{2\dot{p}(1-\dot{p})}$$



The **total impurity** for a data partitioning is just the weighted sum of each partition's impurity

$$\text{Imp}(\{D_1, \dots, D_l\}) = \sum_{i=1}^l \frac{|D_i|}{|D|} \text{Imp}(D_i)$$

Impurity functions for $k > 2$

For more than two classes, the **impurity functions** are defined by the sum of the **per-class impurities** based on “**one versus rest**”

k -class Entropy

$$\text{Imp}(\dot{p}_1, \dots, \dot{p}_k) = \sum_{i=1}^k -\dot{p}_i \log_2(\dot{p}_i) \quad \text{where} \quad \dot{p}_i = \frac{C_i}{\sum_{i=1}^k C_i}$$

k -class Gini index

$$\text{Imp}(\dot{p}) = \sum_{i=1}^k \dot{p}_i (1 - \dot{p}_i)$$

To split a parent node D into children D_1, \dots, D_L we can consider the **purity gain** = $\text{Imp}(D) - \text{Imp}(\{D_1, \dots, D_L\})$

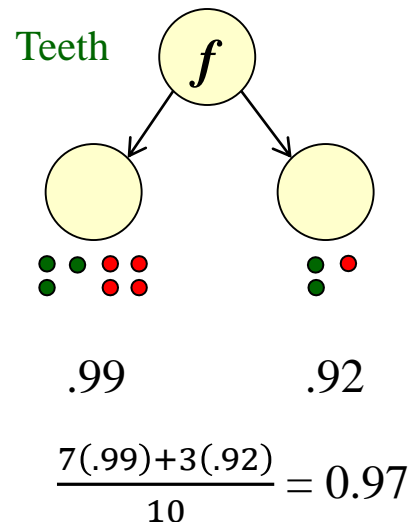
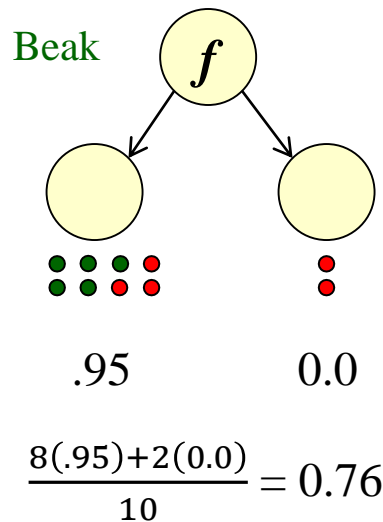
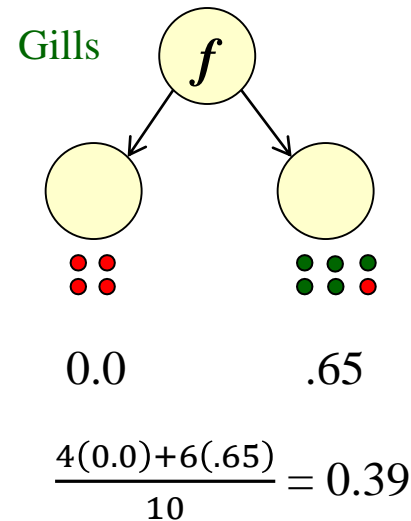
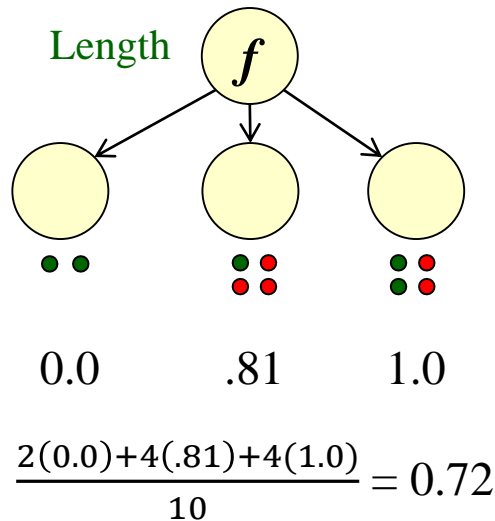
Reminder: What is k ? What is L ?

The number of classes

The number of values (literals) for a given feature F_i

Impurity example

D 



Using the **entropy** measure

$$\text{Imp}(\dot{p}) = -\dot{p} \log_2(\dot{p}) - (1-\dot{p}) \log_2(1-\dot{p})$$

$$\text{Imp}(\{D_1, \dots, D_l\}) = \sum_{i=1}^l \frac{|D_i|}{|D|} \text{Imp}(D_i)$$

Which of these is the best feature to use?

This is the Gills feature in our dolphin example

DT for dolphin example

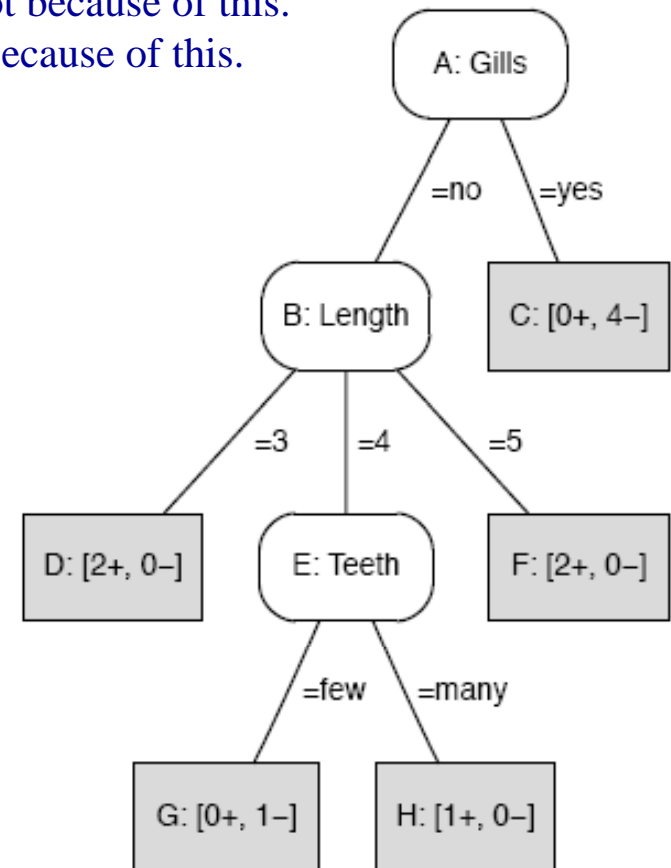
Training data:

1. Length = 3 \wedge Gills = no \wedge Beak = yes \wedge Teeth = many
2. Length = 4 \wedge Gills = no \wedge Beak = yes \wedge Teeth = many
3. Length = 3 \wedge Gills = no \wedge Beak = yes \wedge Teeth = few
4. Length = 5 \wedge Gills = no \wedge Beak = yes \wedge Teeth = many
5. Length = 5 \wedge Gills = no \wedge Beak = yes \wedge Teeth = few
6. Length = 5 \wedge Gills = yes \wedge Beak = yes \wedge Teeth = many
7. Length = 4 \wedge Gills = yes \wedge Beak = yes \wedge Teeth = many
8. Length = 5 \wedge Gills = yes \wedge Beak = no \wedge Teeth = many
9. Length = 4 \wedge Gills = yes \wedge Beak = no \wedge Teeth = many
10. Length = 4 \wedge Gills = no \wedge Beak = yes \wedge Teeth = few

Choose the best feature based on
minimizing impurity of the remaining data

$$\text{Imp}(\dot{p}) = -\dot{p} \log_2(\dot{p}) - (1-\dot{p}) \log_2(1-\dot{p})$$

Why is Gills good?
Not because of this.
Because of this.



DT approach

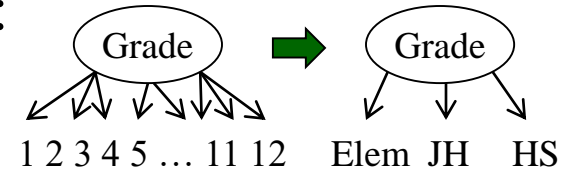
- We've described a **greedy algorithm** – it maximizes each individual choice, but it does not guarantee a global maximum
 - For this we would need the ability to backtrack and reconsider choices based on the **total impurity**

$$\text{Imp}(\{D_1, \dots, D_l\}) = \sum_{i=1}^l \frac{|D_i|}{|D|} \text{Imp}(D_i)$$

- However, it works rather well in practice!
- We can modify the strategy slightly to deal with “messy” (non-separable) data and to limit the size of the tree
 - By not requiring a homogeneous data partition before stopping and assigning a label – i.e., the **Homogeneous(D)** function
 - E.g., if we have a feature separates as **{1000+, 3–}**, that may be good enough – no need to keep checking additional features

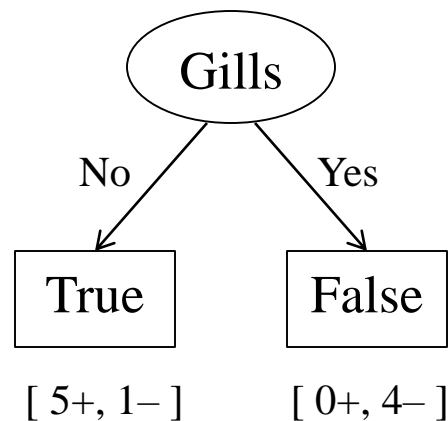
Simplifying decision trees

- Some ways to create simpler decision trees:
 - Merge feature labels and test the difference
 - Enforce a **depth limit** (maximum depth of d)
 - Enforce a **purity threshold** – e.g., if the impurity of a node is $< \epsilon$, turn the node into a leaf (don't expand it further)
 - Enforce a **purity increment threshold** – e.g., if a node expansion increases purity by less than δ , delete the expansion and turn the node into a leaf
 - Build the complete tree and then iteratively **merge leaves** based on lowest purity decrease up to a number of leaves N or a purity δ
 - **Combinations** of depth and purity measures
- If you took 165A, this should remind you of **search** strategies!
 - In fact, we've been discussing various ways of **searching the hypothesis space** to come up with a “good” hypothesis based on our data



Simplifying decision trees

We could simplify the **dolphin** decision tree to this:

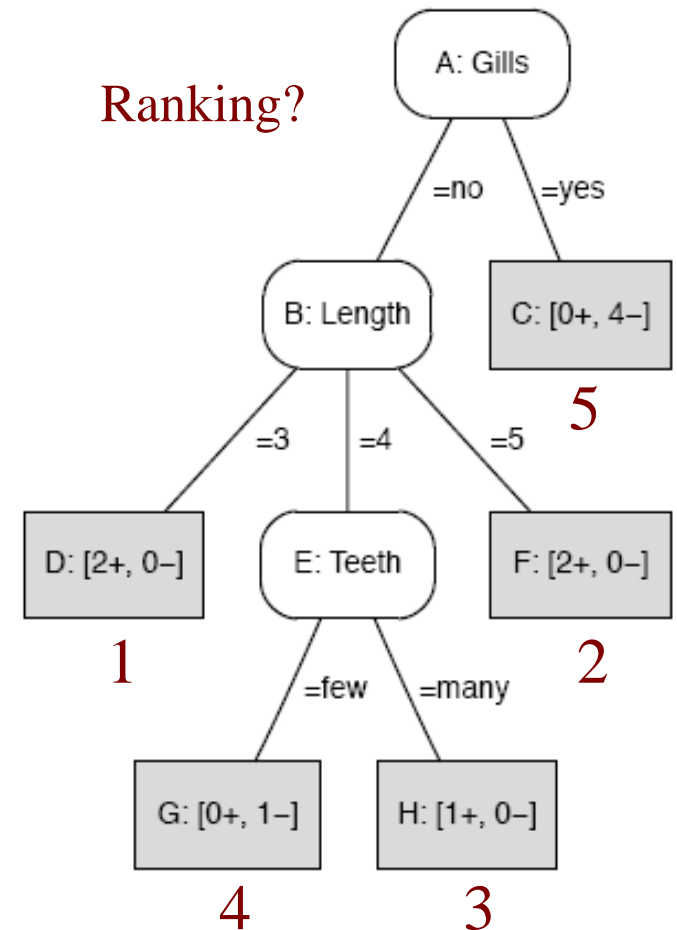


Does this generalize well?

Ranking trees

Note: We can't tell the ranking from the tree structure; only from leaves and their data

- Since a decision tree divides the instance space into segments and we have data for each segment, we can turn the DT into a **ranking model** by evaluating and ordering the segments
- As before, we use empirical probabilities \hat{p} – for segments i and j , order $i > j$ if $\hat{p}_i > \hat{p}_j$
 - May use Laplace correction or m-estimate for smoothing
- As before, we can compute ranking **error rate** and **accuracy**



(This would be a better example if it had more data and non-homogeneous leaves!)

Ranking trees

Dolphin Ranking: D F H G C

- Ranking is with respect to a particular class (e.g., *dolphin*)
- A ranking is on a set of m instances $X = \{ x_1, \dots, x_m \}$
- A decision tree with N leaves will have N different ranks
 - Each instance will have one of those ranks, 1.. N
 - So there are likely to be many ties if N is small or m is large

- | | Leaf |
|--|------|
| 1. Length = 3 \wedge Gills = no \wedge Beak = yes \wedge Teeth = many | D |
| 2. Length = 4 \wedge Gills = no \wedge Beak = yes \wedge Teeth = many | H |
| 3. Length = 3 \wedge Gills = no \wedge Beak = yes \wedge Teeth = few | D |
| 4. Length = 5 \wedge Gills = no \wedge Beak = yes \wedge Teeth = many | F |
| 5. Length = 5 \wedge Gills = no \wedge Beak = yes \wedge Teeth = few | F |
| 6. Length = 5 \wedge Gills = yes \wedge Beak = yes \wedge Teeth = many | C |
| 7. Length = 4 \wedge Gills = yes \wedge Beak = yes \wedge Teeth = many | C |
| 8. Length = 5 \wedge Gills = yes \wedge Beak = no \wedge Teeth = many | C |
| 9. Length = 4 \wedge Gills = yes \wedge Beak = no \wedge Teeth = many | C |
| 10. Length = 4 \wedge Gills = no \wedge Beak = yes \wedge Teeth = few | G |

Leaf

D

H

D

F

F

C

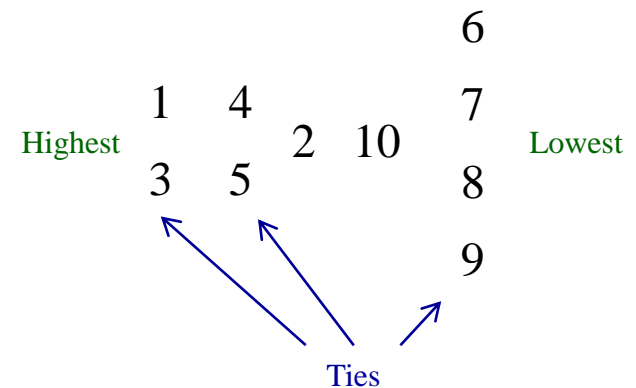
C

C

C

G

From the leaf rankings on the previous slide, the ranking of the 10 instances is:



For this, **accuracy** = 100%

Probability estimation trees

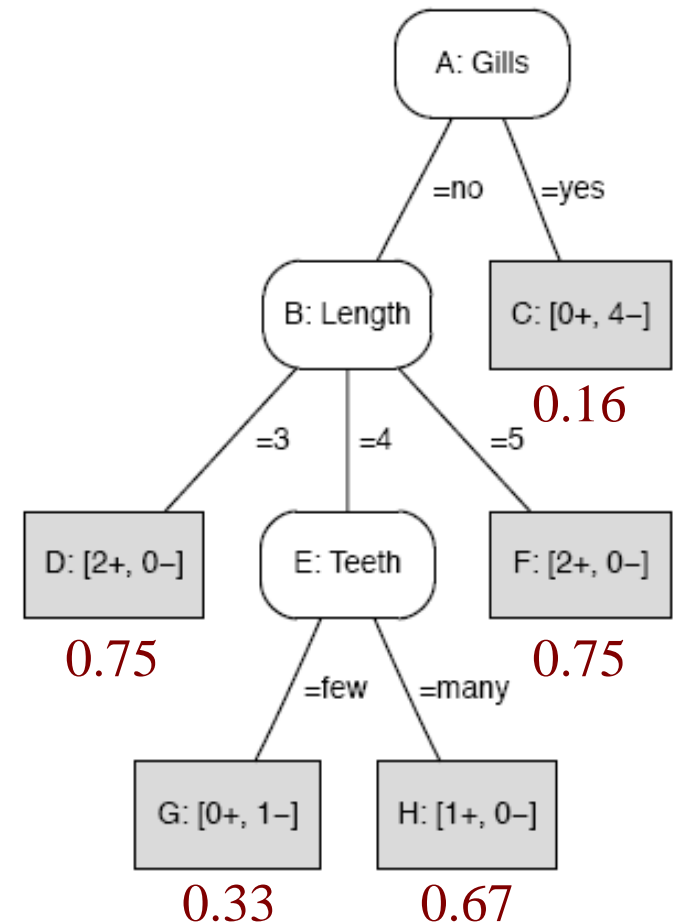
- We can use those empirical probabilities (for each class) calculated for every leaf to create a probability estimation tree
 - A probability classifier, a.k.a. probability estimation
 - Since it's a 2-class problem, we can just show the probability for *dolphin*
- Using Laplace correction, what are the leaf probability estimates?

$$\text{Laplace correction} = \frac{N_i + 1}{|S| + k}$$

$P(\text{dolphin}) = \{ 0.75, 0.75, 0.67, 0.33, 0.16 \}$

Actually showing $P(\text{dolphin}=\text{true} \mid \text{leaf})$ or $P(\text{hypothesis} \mid \text{data})$

$P(\neg \text{dolphin}) = \{ 0.25, 0.25, 0.33, 0.67, 0.84 \}$



Logical models – summary

- In **concept learning** and **decision trees**, we've mostly been discussing **logical models**, based on simple predicate (or first-order) logic
- Pros:
 - English-readable data
 - Intuitive representation and models for people to comprehend
 - Good for explaining the decision-making process
 - Some errors are obvious, easily found and debugged
 - Well suited for some problems
- Cons:
 - Not a good fit for massive amounts of data, for numerical data, for subtle concepts, for things that are difficult to articulate in language
 - I.e., for many of the most important problems ML is being applied to these days!

Where we're going from here

- From **logical models** back to **geometry models** and then on to **probabilistic models**
- There are many uses of logical models, especially decision trees, in machine learning applications
 - DTs are used in many current, practical machine learning systems
- But the focus for some time has moved toward methods that can crunch large amounts of numbers – more and more to **statistical and probabilistic models and methods**
- We'll continue to focus on **classification** and **regression**, as well as **clustering**, and we'll address these problems with a variety of “modern ML” tools and techniques
- Building a linear classifier may seem like a long way from creating intelligent machines that learn and think – but it's not as far as you may think!