

Machine Learning

CS 165B

Prof. Matthew Turk

Monday, May 16, 2016

**T
o
d
a
y**

- Distance metrics and clustering (Ch. 8)

Notes

- HW#4 due on Friday (May 20)

k -Nearest neighbor (k NN) classifiers

- In some cases, the *k -nearest neighbor* method is preferable:
 - Classify a new instance by taking a **vote** of the $k \geq 1$ **nearest exemplars**
 - E.g., in a binary classifier, with $k = 7$, for a new input point the 7 nearest neighbors may include 5 positives and 2 negatives, so we choose positive as the classification
- Or, instead of using a fixed k , vote among all neighbors within a fixed **radius** r
- Or, combine the two, stopping when ($count > k$) or ($dist. > r$)
- May also use **distance weighting** – the closer an exemplar is to the instance, the more its vote counts (e.g., $w_i = \frac{1}{D(x, x_i)}$)
- What about **ties**?
 - Preference to the 1NN
 - Random choice
 - Etc.

Nearest neighbor classification – summary

- NN classifiers are very **fast** to train – $O(n)$ time
 - $n = \#$ of training samples
- But its classification is relatively **slow** – also $O(n)$ time
 - Need to compare the input instance with every stored training example
- Bottom line: **nearest neighbor classifiers** are simple, intuitive, and train quickly
 - But they can be inefficient, may require a good deal of storage, and can't easily represent a specific boundary geometry
- Importantly, NN methods rely on a useful **distance metric**
 - *Nearest* in Euclidian distance, Manhattan distance, Mahalanobis distance, or what?
 - This is problem-dependent
 - **Distance-based** methods


Clustering vs. classification

- Classification vs. clustering
 - In a **classifier**, possible class labels are provided
 - { dog, cat, elephant, mouse, ... }, { spam, ham }, etc.
 - Given in the training data (for supervised classification)
 - In a **clustering** problem, possible labels are the cluster labels learned from the training set
 - { cluster 1, cluster 2, cluster 3, ... }
 - Not given in the training data
- Terminology: In both cases, people often refer to the assigning of labels or clusters to data points (during the learning/training process, or afterwards in testing) as **classification**
 - Even if it's a **clustering** problem!




Clustering

- The goal of clustering is to find clusters (groupings) that are **compact** with respect to the distance metric
- What do we mean by **compactness**?

Is this cluster compact?



It depends....



So we'd like to have a measure of **within-cluster** and **between-cluster** distribution or **scatter**

Scatter matrix

This is a widely-used concept in ML!

From the 4-25 lecture notes:

$$\left[\begin{array}{l} \text{Sample covariance: } \hat{\Sigma}_{ij} = \frac{1}{k} \sum_k (x_{ik} - \hat{\mu}_i)(x_{jk} - \hat{\mu}_j) = \frac{1}{k} S_{ij} \\ \\ \text{If } X \text{ is a matrix that holds all the zero-centered} \\ \text{samples as } \underline{\text{column}} \text{ vectors, then} \end{array} \right. \quad \hat{\Sigma} = \frac{1}{k} \boxed{X_z X_z^T} = \frac{1}{k} S$$

S is the
Scatter matrix

Alternatively, if X_z is a matrix that holds all the
zero-centered samples as row vectors, then

$$S = \boxed{X_z^T X_z}$$

It depends on how we define X_z !

For the **scatter matrix** (and thus the covariance matrix), X_z is **zero-mean**

- That is, the mean data point $\bar{\mathbf{x}}$ (or $\boldsymbol{\mu}$ or $\boldsymbol{\mu}_x$) is first subtracted from every data point \mathbf{x}_i

By the way, the **Gram matrix** is not zero-mean...

Scatter matrix

- If the data D is partitioned into K subsets $\{D_1, D_2, \dots, D_K\}$ then the scatter matrix can be written as

$$S = \left(\sum_{j=1}^K S_j \right) + B$$


The diagram illustrates the decomposition of the scatter matrix S . The term $\sum_{j=1}^K S_j$ is identified as the sum of subset scatter matrices, which are also referred to as within-cluster scatter matrices. The term B is identified as the scatter matrix of the partition means, also referred to as the between-cluster scatter matrix. Blue arrows point from the text labels to their respective terms in the equation.

Subset scatter matrices
Within-cluster scatter matrices

Scatter matrix of the partition means
Between-cluster scatter matrix

To compute B , replace every point in D with the mean of its partition D_i and compute the scatter matrix

Example: Scatter matrix



$$\mathbf{X} = [\mathbf{x}_1 \quad \mathbf{x}_2 \quad \mathbf{x}_3 \quad \mathbf{x}_4 \quad \mathbf{x}_5] = \begin{bmatrix} 2 & 4 & 0 & 3 & 5 \\ 0 & 1 & 4 & 4 & 2 \end{bmatrix}$$

The **Gram matrix** is ...

$$\mathbf{G} = \mathbf{X}^T \mathbf{X} = \begin{bmatrix} 4 & 8 & 0 & 6 & 10 \\ 8 & 17 & 4 & 16 & 22 \\ 0 & 4 & 16 & 16 & 8 \\ 6 & 16 & 16 & 25 & 23 \\ 10 & 22 & 8 & 23 & 29 \end{bmatrix}$$

($k \times k$), where k is the number of data points

$$\bar{\mathbf{x}} = \frac{1}{5} \sum_{i=1}^5 \mathbf{x}_i = \frac{1}{5} \begin{bmatrix} 14 \\ 11 \end{bmatrix} = \begin{bmatrix} 2.8 \\ 2.2 \end{bmatrix}$$

$$\mathbf{X}_z = [\mathbf{x}_1 - \bar{\mathbf{x}} \quad \mathbf{x}_2 - \bar{\mathbf{x}} \quad \mathbf{x}_3 - \bar{\mathbf{x}} \quad \mathbf{x}_4 - \bar{\mathbf{x}} \quad \mathbf{x}_5 - \bar{\mathbf{x}}] = \begin{bmatrix} -0.8 & 1.2 & -2.8 & 0.2 & 2.2 \\ -2.2 & -1.2 & 1.8 & 1.8 & -0.2 \end{bmatrix}$$


The **Scatter matrix** is ...

$$\mathbf{S} = \mathbf{X}_z \mathbf{X}_z^T = \begin{bmatrix} 14.8 & -4.8 \\ -4.8 & 12.8 \end{bmatrix}$$

($N \times N$), where N is the dimensionality of the data points

Example: Scatter matrix via partitions

$$S = \left(\sum_{j=1}^K S_j \right) + B$$



Partition means

$$\begin{bmatrix} 1.5 & 1.5 & 3.67 & 3.67 & 3.67 \\ 4 & 4 & 1 & 1 & 1 \end{bmatrix} \quad \mu_B = \begin{bmatrix} 2.8 \\ 2.2 \end{bmatrix}$$

Zero-mean partition means

$$B_z = \begin{bmatrix} -1.3 & -1.3 & 13/15 & 13/15 & 13/15 \\ 1.8 & 1.8 & -1.2 & -1.2 & -1.2 \end{bmatrix}$$

Between-cluster scatter matrix

$$B = B_z B_z^T = \begin{bmatrix} 5.633 & -7.8 \\ -7.8 & 10.8 \end{bmatrix}$$

Scatter matrix of D_1

$$S_1 = \left[\mathbf{x}_3 - \begin{bmatrix} 1.5 \\ 4 \end{bmatrix} \quad \mathbf{x}_4 - \begin{bmatrix} 1.5 \\ 4 \end{bmatrix} \right] \left[\mathbf{x}_3 - \begin{bmatrix} 1.5 \\ 4 \end{bmatrix} \quad \mathbf{x}_4 - \begin{bmatrix} 1.5 \\ 4 \end{bmatrix} \right]^T = \begin{bmatrix} -1.5 & 1.5 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} -1.5 & 1.5 \\ 0 & 0 \end{bmatrix}^T = \begin{bmatrix} 4.5 & 0 \\ 0 & 0 \end{bmatrix}$$

Scatter matrix of D_2

$$S_2 = \begin{bmatrix} -5/3 & 1/3 & 4/3 \\ -1 & 0 & 1 \end{bmatrix} \begin{bmatrix} -5/3 & 1/3 & 4/3 \\ -1 & 0 & 1 \end{bmatrix}^T = \begin{bmatrix} 4.67 & 3 \\ 3 & 2 \end{bmatrix}$$

$$S = S_1 + S_2 + B = \begin{bmatrix} 4.5 & 0 \\ 0 & 0 \end{bmatrix} + \begin{bmatrix} 4.67 & 3 \\ 3 & 2 \end{bmatrix} + \begin{bmatrix} 5.633 & -7.8 \\ -7.8 & 10.8 \end{bmatrix} = \begin{bmatrix} 14.8 & -4.8 \\ -4.8 & 12.8 \end{bmatrix}$$

Scatter

- The **scatter** of \mathbf{X} is defined as the **trace** of the scatter matrix
 - The *trace* is the sum of the diagonal elements of a square matrix

$$\begin{aligned}\text{Scat}(\mathbf{X}) &= \text{Tr}(\mathbf{S}) = \text{Tr}(\mathbf{X}_z \mathbf{X}_z^T) \\ &= \text{Tr} \left(\begin{bmatrix} 14.8 & -4.8 \\ -4.8 & 12.8 \end{bmatrix} \right) = 14.8 + 12.8 = 27.6\end{aligned}$$

- Since \mathbf{S} can be decomposed into partitions, so can $\text{Scat}(\mathbf{X})$

$$\text{Scat}(D) = \underbrace{\sum_{j=1}^K \text{Scat}(D_j)}_{\substack{\uparrow \\ \text{Want to choose} \\ \text{partitions that} \\ \text{minimize this}}} + \underbrace{\sum_{j=1}^K |D_j| \|\mu_j - \mu\|^2}_{\substack{\uparrow \\ \text{Equivalent to} \\ \text{maximizing this}}}$$

Fixed for a given data set

This is the goal of
k-means clustering

K-means clustering

- The general **K-means clustering** problem is **NP-complete**, so there is no efficient solution to find the **optimal clustering** (data partition)
- A widely-used **heuristic** algorithm for clustering is also known as the **K-means algorithm**, but it is not optimal
 - It will converge to a solution, but there is no guarantee that the solution is the best one (the global minimum of scatter)
 - But it works quite well in most cases!
- Typically, the **K-means algorithm** would be run several times (with a random starting point) and then the best solution is selected
 - I.e., the solution with the smallest **within-cluster scatter**

K-means algorithm

K is an input parameter

Algorithm $KMeans(D, K)$ – K -means clustering using Euclidean distance Dis_2 .

Input : data $D \subseteq \mathbb{R}^d$; number of clusters $K \in \mathbb{N}$.

Output : K cluster means $\mu_1, \dots, \mu_K \in \mathbb{R}^d$.

randomly initialise K vectors $\mu_1, \dots, \mu_K \in \mathbb{R}^d$;

repeat

 assign each $\mathbf{x} \in D$ to $\operatorname{argmin}_j Dis_2(\mathbf{x}, \mu_j)$; \leftarrow *1-Nearest neighbor assignment*

for $j = 1$ to K **do**

$D_j \leftarrow \{\mathbf{x} \in D \mid \mathbf{x} \text{ assigned to cluster } j\}$; \leftarrow *Partition defined by assignment*


$\mu_j = \frac{1}{|D_j|} \sum_{\mathbf{x} \in D_j} \mathbf{x}$; \leftarrow *Re-compute the cluster mean*

end


until no change in μ_1, \dots, μ_K ;

return μ_1, \dots, μ_K ;


1




2



3



4




K-medoids algorithm

- In some problems, the cluster **exemplars** (representatives) are required to be **data points**
 - As opposed to using the **mean** of the cluster points, for example, since the mean is most likely not a point in the data set
- The concept of **medoid** is useful here – the **medoid** of a set of points is the point with the **minimal average dissimilarity** (distance) to all other points in the set
 - Using some distance metric: Euclidian, L1, etc.
 - This is a generalization of the concept of **median** to multiple dimensions
- K-means can be modified to use **data points** as exemplars rather than **means**


K-medoids algorithm

Cluster mean



Location that minimizes the sum of squared distances to points

Cluster medoid



Point that minimizes the sum of squared distances to points

K-medoids algorithm

Algorithm $\text{KMedoids}(D, K, \text{Dis})$ – K -medoids clustering using arbitrary distance metric Dis .

Input : data $D \subseteq \mathcal{X}$; number of clusters $K \in \mathbb{N}$;
distance metric $\text{Dis} : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$.

Output : K medoids $\mu_1, \dots, \mu_K \in D$, representing a predictive clustering of \mathcal{X} .
randomly pick K data points $\mu_1, \dots, \mu_K \in D$;

repeat

 assign each $\mathbf{x} \in D$ to $\arg\min_j \text{Dis}(\mathbf{x}, \mu_j)$;

for $j = 1$ to K **do**

$D_j \leftarrow \{\mathbf{x} \in D \mid \mathbf{x} \text{ assigned to cluster } j\}$;

$\mu_j = \arg\min_{\mathbf{x} \in D_j} \sum_{\mathbf{x}' \in D_j} \text{Dis}(\mathbf{x}, \mathbf{x}')$; \leftarrow *Re-compute the cluster medoid*

end

until no change in μ_1, \dots, μ_K ;

return μ_1, \dots, μ_K ;

Kernel K-means clustering

Algorithm $\text{Kernel-KMeans}(D, K)$ – K -means clustering using kernelised distance Dis_{κ} .

Input : data $D \subseteq \mathcal{X}$; number of clusters $K \in \mathbb{N}$.

Output : K -fold partition $D_1 \uplus \dots \uplus D_K = D$.

randomly initialise K clusters D_1, \dots, D_K ;

repeat

 assign each $\mathbf{x} \in D$ to $\arg\min_j \frac{1}{|D_j|} \sum_{\mathbf{y} \in D_j} \text{Dis}_{\kappa}(\mathbf{x}, \mathbf{y})$; \leftarrow *Re-assign each point to a partition according to the minimum average (kernel) distance*

for $j = 1$ to K **do**

$D_j \leftarrow \{\mathbf{x} \in D \mid \mathbf{x} \text{ assigned to cluster } j\}$;

end

until no change in D_1, \dots, D_K ;

return D_1, \dots, D_K ;

As before, replace the dot product with a kernel function κ


$$\text{Dis}_2(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\|_2 = \sqrt{(\mathbf{x} - \mathbf{y}) \cdot (\mathbf{x} - \mathbf{y})} = \sqrt{\mathbf{x} \cdot \mathbf{x} - 2\mathbf{x} \cdot \mathbf{y} + \mathbf{y} \cdot \mathbf{y}}$$

$$\text{Dis}_{\kappa}(\mathbf{x}, \mathbf{y}) = \sqrt{\kappa(\mathbf{x}, \mathbf{x}) - 2\kappa(\mathbf{x}, \mathbf{y}) + \kappa(\mathbf{y}, \mathbf{y})}$$


Clustering

- These clustering methods are **distance-based** methods that do not take into account information about the cluster **shape**
 - This can lead to counter-intuitive and unwanted results

$K = 2$



K-means finds the
two natural clusters



After scaling the data in
y, the found clusters are
not what we expected!

Summary: Distance methods and clustering

- Euclidian distance may not always be the right choice
- **Similarity** is a function of **distance**
- **Nearest neighbor** methods assign classes/clusters based on distances to points or **exemplars**, not based on computed boundaries
- For good clustering, we want high **within-class** (intra-class) similarity and low **between-class** (inter-class) similarity
- The **scatter matrix** is an important structure in clustering
- The K-means algorithm (and variations) is widely used