

Manhattan (L1) distance:  $d(x, y) = \sum_{i=1}^d |x_i - y_i|$

Euclidian (L2) distance:  $d(x, y) = \|x - y\| = \left( \sum_{i=1}^d (x_i - y_i)^2 \right)^{1/2}$

Minkowski (Lp) distance:  $d(x, y) = \left( \sum_{i=1}^d |x_i - y_i|^p \right)^{1/p}$

$$\text{Laplace correction} = \frac{N_i + 1}{|S| + k} \qquad \text{m-estimate} = \frac{N_i + m\pi_i}{|S| + m}$$

---

**Algorithm** *GrowTree(D, F)* – grow a feature tree from training data.

---

**Input** : data  $D$ ; set of features  $F$ .

**Output** : feature tree  $T$  with labelled leaves.

**if** *Homogeneous(D)* **then return** *Label(D)*;

$S \leftarrow \text{BestSplit}(D, F)$  ; // e.g., *BestSplit-Class* (Algorithm 5.2)

split  $D$  into subsets  $D_i$  according to the literals in  $S$ ;

**for** each  $i$  **do**

**if**  $D_i \neq \emptyset$  **then**  $T_i \leftarrow \text{GrowTree}(D_i, F)$  ;

**else**  $T_i$  is a leaf labelled with *Label(D)*;

**end**

**return** a tree whose root is labelled with  $S$  and whose children are  $T_i$

---

### Impurity measures:

Minority class

$$\text{Imp}(\dot{p}) = \min(\dot{p}, 1 - \dot{p})$$

Gini index

$$\text{Imp}(\dot{p}) = 2\dot{p}(1 - \dot{p})$$

Entropy

$$\text{Imp}(\dot{p}) = -\dot{p} \log_2(\dot{p}) - (1 - \dot{p}) \log_2(1 - \dot{p})$$

$\sqrt{\text{Gini}}$  index

$$\text{Imp}(\dot{p}) = \sqrt{2\dot{p}(1 - \dot{p})}$$

### Total impurity:

$$\text{Imp}(\{D_1, \dots, D_l\}) = \sum_{i=1}^l \frac{|D_i|}{|D|} \text{Imp}(D_i)$$

Bayes Rule:

$$P(H_i | D) = \frac{P(D | H_i) P(H_i)}{P(D)}$$

$$\text{False positive rate (FPR)} = \frac{FP}{N} = \alpha$$

$$\text{Accuracy} = \frac{TP + TN}{P + N} = \left(\frac{P}{P + N}\right) TPR + \left(\frac{N}{P + N}\right) TNR$$

$$\text{False negative rate (FNR)} = \frac{FN}{P} = \beta$$

$$\text{Error rate} = \frac{FP + FN}{P + N}$$

$$\text{True positive rate (TPR)} = \frac{TP}{P}$$

$$\text{Precision} = \frac{TP}{\hat{P}}$$

$$\text{True negative rate (TNR)} = \frac{TN}{N}$$

$$\text{Accuracy} + \text{error rate} = 1$$

$$\text{Average recall} = \frac{TPR + TNR}{2}$$

$$\text{Ranking classifier error rate: } rank\text{-err} = err / PN$$

$$\text{Ranking classifier accuracy: } rank\text{-acc} = 1 - rank\text{-err}$$

Multivariate least-squares regression  
(homogeneous representation):

$$\mathbf{y} = \mathbf{X}\mathbf{w} + \epsilon$$

$$\begin{aligned}\hat{\mathbf{w}} &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \\ &= \mathbf{S}^{-1} \mathbf{X}^T \mathbf{y}\end{aligned}$$

Classifier margin for point  $\mathbf{x}$

$$\mathbf{z}(\mathbf{x}) = \frac{y(\mathbf{w}^T \mathbf{x} - t)}{\|\mathbf{w}\|} = \frac{m}{\|\mathbf{w}\|}$$

Non-homogeneous  
representation

PAC learning outputs, with probability at least  $1 - \delta$ , a hypothesis  $h$  such that  $err_D < \epsilon$

---

**Algorithm**  $\text{Perceptron}(D, \eta)$  – train a perceptron for linear classification.

---

**Input** : labelled training data  $D$  in homogeneous coordinates; learning rate  $\eta$ .

**Output** : weight vector  $\mathbf{w}$  defining classifier  $\hat{y} = \text{sign}(\mathbf{w} \cdot \mathbf{x})$ .

$\mathbf{w} \leftarrow \mathbf{0}$  ; // Other initialisations of the weight vector are possible

$\text{converged} \leftarrow \text{false}$ ;

**while**  $\text{converged} = \text{false}$  **do**

$\text{converged} \leftarrow \text{true}$ ;

**for**  $i = 1$  to  $|D|$  **do**

**if**  $y_i \mathbf{w} \cdot \mathbf{x}_i \leq 0$  // i.e.,  $\hat{y}_i \neq y_i$

**then**

$\mathbf{w} \leftarrow \mathbf{w} + \eta y_i \mathbf{x}_i$ ;

$\text{converged} \leftarrow \text{false}$ ; // We changed  $\mathbf{w}$  so haven't converged yet

**end**

**end**

**end**

---



---

**Algorithm**  $\text{DualPerceptron}(D)$  – perceptron training in dual form.

---

**Input** : labelled training data  $D$  in homogeneous coordinates.

**Output** : coefficients  $\alpha_i$  defining weight vector  $\mathbf{w} = \sum_{i=1}^{|D|} \alpha_i y_i \mathbf{x}_i$ .

$\alpha_i \leftarrow 0$  for  $1 \leq i \leq |D|$ ;

$\text{converged} \leftarrow \text{false}$ ;

**while**  $\text{converged} = \text{false}$  **do**

$\text{converged} \leftarrow \text{true}$ ;

**for**  $i = 1$  to  $|D|$  **do**

**if**  $y_i \sum_{j=1}^{|D|} \alpha_j y_j \mathbf{x}_i \cdot \mathbf{x}_j \leq 0$  **then**

$\alpha_i \leftarrow \alpha_i + 1$ ;

$\text{converged} \leftarrow \text{false}$ ;

**end**

**end**

**end**

---

Sample covariance:  $\hat{\Sigma}_{ij} = \frac{1}{k} \sum_k (x_{ik} - \hat{\mu}_i)(x_{jk} - \hat{\mu}_j) = \frac{1}{k} S_{ij}$

If  $X$  is a matrix that holds all the zero-centered samples as column vectors, then

$$\hat{\Sigma} = \frac{1}{k} \mathbf{X} \mathbf{X}^T = \frac{1}{k} \mathbf{S}$$

$\mathbf{S}$  is the scatter matrix

If  $X$  is not zero-centered, then

$$\mathbf{G} = \mathbf{X}^T \mathbf{X}$$

$\mathbf{G}$  is the Gram matrix

Soft margin optimization problem:

$$\mathbf{w}^*, t^*, \xi_i^* = \arg \min_{\mathbf{w}, t, \xi_i} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$$

subject to  $y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1 - \xi_i$  and  $\xi_i \geq 0, 1 \leq i \leq n$

Chebyshev distance:

$$L_\infty(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\|_\infty = \max_i |x_i - y_i|$$

Hamming distance:

$$L_0(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\|_0 = \text{count}(|x_i - y_i| > 0)$$

Mahalanobis distance:

$$D_M(\mathbf{x}, \mathbf{y}) = \sqrt{(\mathbf{x} - \mathbf{y})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \mathbf{y})}$$

---

**Algorithm**  $K\text{Means}(D, K)$  –  $K$ -means clustering using Euclidean distance  $\text{Dis}_2$ .

---

**Input** : data  $D \subseteq \mathbb{R}^d$ ; number of clusters  $K \in \mathbb{N}$ .

**Output** :  $K$  cluster means  $\mu_1, \dots, \mu_K \in \mathbb{R}^d$ .

randomly initialise  $K$  vectors  $\mu_1, \dots, \mu_K \in \mathbb{R}^d$ ;

**repeat**

    assign each  $\mathbf{x} \in D$  to  $\arg \min_j \text{Dis}_2(\mathbf{x}, \mu_j)$ ;

**for**  $j = 1$  to  $K$  **do**

$D_j \leftarrow \{\mathbf{x} \in D \mid \mathbf{x} \text{ assigned to cluster } j\}$ ;

$\mu_j = \frac{1}{|D_j|} \sum_{\mathbf{x} \in D_j} \mathbf{x}$ ;

**end**

**until** no change in  $\mu_1, \dots, \mu_K$ ;

**return**  $\mu_1, \dots, \mu_K$ ;

---

---

**Algorithm** Bagging( $D, T, \mathcal{A}$ ) – train an ensemble of models from bootstrap samples.

---

**Input** : data set  $D$ ; ensemble size  $T$ ; learning algorithm  $\mathcal{A}$ .

**Output** : ensemble of models whose predictions are to be combined by voting or averaging.

**for**  $t = 1$  to  $T$  **do**

build a bootstrap sample  $D_t$  from  $D$  by sampling  $|D|$  data points with replacement;  
run  $\mathcal{A}$  on  $D_t$  to produce a model  $M_t$ ;

**end**

**return**  $\{M_t | 1 \leq t \leq T\}$

---



---

**Algorithm** Boosting( $D, T, \mathcal{A}$ ) – train an ensemble of binary classifiers from reweighted training sets.

---

**Input** : data set  $D$ ; ensemble size  $T$ ; learning algorithm  $\mathcal{A}$ .

**Output** : weighted ensemble of models.

$w_{1i} \leftarrow 1/|D|$  for all  $x_i \in D$ ; // start with uniform weights

**for**  $t = 1$  to  $T$  **do**

run  $\mathcal{A}$  on  $D$  with weights  $w_{ti}$  to produce a model  $M_t$ ;  
calculate weighted error  $\epsilon_t$ ;

**if**  $\epsilon_t \geq 1/2$  **then**

set  $T \leftarrow t - 1$  and break

**end**

$\alpha_t \leftarrow \frac{1}{2} \ln \frac{1-\epsilon_t}{\epsilon_t}$ ; // confidence for this model

$w_{(t+1)i} \leftarrow \frac{w_{ti}}{2\epsilon_t}$  for misclassified instances  $x_i \in D$ ; // increase weight

$w_{(t+1)j} \leftarrow \frac{w_{tj}}{2(1-\epsilon_t)}$  for correctly classified instances  $x_j \in D$ ; // decrease

**end**

**return**  $M(x) = \sum_{t=1}^T \alpha_t M_t(x)$

---

Sigmoid function:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Backpropagation error for output units:

$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$

Backpropagation error for hidden units:

$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{kh} \delta_k$$